

Harshil Shah – SEC01 (NUID 002780887)

# Big Data System Engineering with Scala

## Spring 2023

### Assignment - 7



- GitHub Repo URL - <https://github.com/harshilshahneu/CSYE7200-Harshil-Shah>

-Kaggle dataset -

<https://www.kaggle.com/datasets/rounakbanik/the-movies-dataset> (ratings.csv)

---

### - List of Tasks Implemented

ou are required to analyze a movie rating dataset. The data is stored in a CSV file (either use the one in the repository or download the latest from Kaggle). You need to read this file into spark and calculate the mean rating and standard deviation for all movies. There is no test case provided for you, so you need to write your own test cases to ensure that at least your program works well.

You can refer to *WordCount.scala* file for the basic structure. Notice that you need to use specific Spark version 3.2.1 for Scala 2.12.x support. If you can run it using a version with Scala 2.13, then go ahead.

Note also that there is a module in the CSYE7200 repository called *spark-csv*. If you use that, you will have to edit the *build.sbt* file. You can use that code of course to get started with (show your results if you do this). **However**, you also need to read the CSV file using the Spark utility (see <https://spark.apache.org/docs/3.2.1/sql-data-sources-csv.html>) and then create a method that accepts a DataFrame and returns the processed DataFrame.

You need to provide your code (in your own repository) together with the mean/std. dev. Don't forget to say where exactly you got the CSV file from.

---

### - Code

**MovieAnalyzer.scala**

```

1  import org.apache.spark.sql.SparkSession
2  import org.apache.spark.sql.functions.{avg, stddev}
3  import org.apache.spark.sql.types.DoubleType
4
5  object MovieAnalyzer {
6
7      def main(args: Array[String]): Unit = {
8
9          val spark = SparkSession.builder()
10             .appName("MovieAnalyzer")
11             .master("local[*]")
12             .getOrCreate()
13
14          val df = spark.read
15             .option("header", true)
16             .option("inferSchema", true)
17             .csv("spark-movie-rating/src/main/resources/ratings.csv")
18
19          // Calculate the mean rating for each movie
20          val ratings = df.groupBy("movieId")
21             .agg(avg("rating").cast(DoubleType).alias("Mean"),
22                 stddev("rating").cast(DoubleType).alias("Standard Deviation"))
23             .na.fill(0.0, Seq("Standard Deviation"))
24
25          // Show the resulting DataFrame
26          ratings.show()
27          println("count : " + ratings.count()) //45115
28          spark.stop()
29      }
30  }
31

```

**MovieDatabaseAnalyzerTest.scala**

```

1  package edu.neu.coe.csye7200.csv
2
3  import org.apache.spark.sql.{SparkSession}
4  import org.scalatest.flatspec.AnyFlatSpec
5  import org.scalatest.matchers.should.Matchers
6  import org.apache.spark.sql.{DataFrame}
7
8
9  class MovieDatabaseAnalyzerTest extends AnyFlatSpec with Matchers {
10
11    implicit val sparkSession: SparkSession = SparkSession
12      .builder()
13      .appName("MovieDatabaseAnalyzer")
14      .master("local[*]")
15      .getOrCreate()
16
17
18    private var spark: SparkSession = _
19    private var ratings: DataFrame = _
20
21
22    "MovieAnalyzer" should "calculate the mean and standard deviation ratings for each movie" in {
23      // Check that the DataFrame has the expected number of rows
24      assert(ratings.count() == 45115)
25    }
26  }
27

```

## - Results

(only 20 rows from the DF)

```

+-----+-----+-----+
|movieId|      Mean|Standard Deviation|
+-----+-----+-----+
| 1645| 3.516589990241182|0.9682679423119045|
| 1591|2.6416020262782967| 1.100915545891352|
| 3175| 3.586550320670942|0.9813643339255932|
| 1580|3.5733178489322874|0.9056628458885891|
| 68135| 3.097457627118644|1.0896833315419254|
|  471| 3.654817548175482|0.9400255645903338|
| 1088| 3.239810636881426|1.1413956305523005|
| 1959|3.6369782971619364|1.0299294156877061|

```

2122	2.634513274336283	1.1126183541067707
2866	3.6019714479945617	0.7823322374774501
36525	3.4823726916620035	0.9356970501427126
3918	2.9595715272978578	1.1492630792872496
3997	2.077287716405606	1.1282243238913907
6620	3.789404132628544	0.8140886808308354
1238	3.9629796163069546	0.9112289644602707
2142	3.038054538054538	0.9990206728801857
8638	3.9675026123301986	0.9140757179280239
2366	3.4740872335211956	1.114247520550539
1342	2.9637979902087093	1.0790487909275213
3794	3.250574712643678	1.1139737559449223
+-----+-----+-----+

---

Execution Screenshots

```

+-----+-----+-----+
|movieId|          Mean|Standard Deviation|
+-----+-----+-----+
|  1645| 3.516589990241182|0.9682679423119045|
|  1591|2.6416020262782967| 1.100915545891352|
|  3175| 3.586550320670942|0.9813643339255932|
|  1580|3.5733178489322874|0.9056628458885891|
| 68135| 3.097457627118644|1.0896833315419254|
|   471| 3.654817548175482|0.9400255645903338|
|  1088| 3.239810636881426|1.1413956305523005|
|  1959|3.6369782971619364|1.0299294156877061|
|  2122| 2.634513274336283|1.1126183541067707|
|  2866|3.6019714479945617|0.7823322374774501|
| 36525|3.4823726916620035|0.9356970501427126|
|   3918|2.9595715272978578|1.1492630792872496|
|   3997| 2.077287716405606|1.1282243238913907|
|   6620| 3.789404132628544|0.8140886808308354|
|   1238|3.9629796163069546|0.9112289644602707|
|   2142| 3.038054538054538|0.9990206728801857|
|   8638|3.9675026123301986|0.9140757179280239|
|   2366|3.4740872335211956| 1.114247520550539|
|   1342|2.9637979902087093|1.0790487909275213|
|   3794| 3.250574712643678|1.1139737559449223|
+-----+-----+-----+
only showing top 20 rows

```

▶

✓

⊘

↓<sup>a</sup><sub>z</sub>

↓≡

≡

÷

↑

↓

🕒

>>

▶

🔄

🔧

■

📷

📺

>>

▼ ✓ Test Results 3 sec 452 ms

▼ ✓ MovieDatabaseAnalyzer 3 sec 452 ms

▼ ✓ parseResource 3 sec 452 ms

✓ should get movie 3 sec 452 ms